

# Practical Machine Learning Project

Antonio Rubio Calzado

29 de mayo de 2017

## Introduction.

For this project we are given data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. The training data is labeled with the manner in which they did the exercise. We are asked to develop an algorithm able to predict this label and perform it on a testing set (data without this labels).

## Cleaning and processing data.

Firstly, we download the training and testing data on our computer respectively from:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

If we open our raw data, there are multiple fields with "NA", "#DIV/0!" or "", so we will think they are NA values when reading it.

```
setwd("C:/Users/arubioca/Desktop/MLPROJECT")
train <- read.csv("pml-training.csv",dec=".",na.strings = c("NA","",
"#DIV/0!"))
test <- read.csv("pml-testing.csv",dec=".",na.strings = c("NA","",
"#DIV/0!"))
```

Now, let's load the libraries necessary for this project

```
suppressMessages(library(caret))
suppressMessages(library(randomForest))
suppressMessages(library(rpart))
```

The next step is using the **caret** package to create a partition of our training data in two subgroups: The first one will have the 60% of the size of this data and the other one the remaining 40%. This will be used to do cross-validations of our future predictive algorithms.

```
set.seed(100)
inTrain <- createDataPartition(train$classe, p = 0.6, list = FALSE)
training <- train[inTrain,]
testing <- train[-inTrain,]
```

Let's remove near zero variance predictors from our data. With the following commands, we are detecting the variables that are near zero variance predictors in the train dataset.

```
set.seed(100)
nzv <- nearZeroVar(train, saveMetrics = TRUE)
rm_nzv_index <- which(nzv$zeroVar == TRUE)
```

Also notice that many fields in our test.csv are constantly NA. We this procedure, we save in a vector the columns of the train dataset with this behaviour:

```
rm_NA_test_index <- c()
for (i in 1:160){
  if (all(is.na(test[,i]))){
    rm_NA_test_index <- c(rm_NA_test_index,i)
  }
}
```

Now we are going to create a vector with the variables that can be dropped from our train dataset, since they are constantly NA or near zero variance variables.

```
index <- unique(c(rm_nzv_index, rm_NA_test_index))
training_new <- training[,-index]
testing_new <- testing[,-index]
```

The variables X, user\_name, raw\_timestamp\_part\_1, raw\_timestamp\_part\_2, cvtd\_timestamp, new\_window and num\_window can be omitted in our study, so we also dropped them:

```
training_new <- training_new[,-c(1,2,3,4,5,6,7)]
testing_new <- testing_new[,-c(1,2,3,4,5,6,7)]
```

## Predictive Models

We start fitting a CART-tree model on our training set, using the package rpart:

```
set.seed(100)
model_tree <- rpart(classe ~ ., data = training_new, method = "class")
```

After training phase, we are studying the out-of-sample error of this model, by predicting a classe label on our testing set and comparing those predictions with the real classe labels via the confusion Matrix:

```
predictions_tree <- predict(model_tree, testing_new, type = "class")
confusionMatrix(predictions_tree, testing_new$classe)
```

```
## Confusion Matrix and Statistics
```

```
##
##              Reference
## Prediction    A    B    C    D    E
##      A 1968  353   20  121   41
##      B   48  845   83   40   92
##      C   58  143 1170  192  163
##      D   85  111   81  837   74
##      E   73   66   14   96 1072
```


```
##
## Overall Statistics
##
##           Accuracy : 0.751
##           95% CI : (0.7412, 0.7605)
##       No Information Rate : 0.2845
##       P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6839
##  McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.8817  0.5567  0.8553  0.6509  0.7434
## Specificity      0.9047  0.9584  0.9142  0.9465  0.9611
## Pos Pred Value   0.7863  0.7626  0.6779  0.7045  0.8115
## Neg Pred Value   0.9506  0.9001  0.9676  0.9326  0.9433
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2508  0.1077  0.1491  0.1067  0.1366
## Detection Prevalence 0.3190 0.1412 0.2200 0.1514 0.1684
## Balanced Accuracy 0.8932 0.7575 0.8847 0.7987 0.8523
```

The accuracy of the model is only 75.1%, so it's a good idea try to fit another different model.

Let's repeat this proccess with a random forest algorithm. We start training the model using the `randomForest` package:

```
set.seed(100)
model_rf <- randomForest(classe ~ ., data = training_new)
```

The following plot shows a graph of Error vs Trees in the previous random forest model:



Now we are checking the accuracy of this model:

```
set.seed(100)
predictions_rf <- predict(model_rf, testing_new, type = "class")
confusionMatrix(predictions_rf, testing_new$classe)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction    A    B    C    D    E
```

```
##           A 2230     6     1     0     0
```

```
##           B   2 1509     8     0     0
```

```
##           C    0     3 1356    15     2
```

```
##           D    0     0     3 1270     6
```

```
##           E    0     0     0     1 1434
```

```
##
```

```
## Overall Statistics
```

```
##
```

```
##           Accuracy : 0.994
```

```
##           95% CI : (0.992, 0.9956)
```

```
##           No Information Rate : 0.2845
```

```
##           P-Value [Acc > NIR] : < 2.2e-16
```

```
##
```

```
##           Kappa : 0.9924
```

```
##           Mcnemar's Test P-Value : NA
```

```
##
```

```
## Statistics by Class:
```

```
##
##               Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9991   0.9941   0.9912   0.9876   0.9945
## Specificity      0.9988   0.9984   0.9969   0.9986   0.9998
## Pos Pred Value   0.9969   0.9934   0.9855   0.9930   0.9993
## Neg Pred Value   0.9996   0.9986   0.9981   0.9976   0.9988
## Prevalence       0.2845   0.1935   0.1744   0.1639   0.1838
## Detection Rate   0.2842   0.1923   0.1728   0.1619   0.1828
## Detection Prevalence 0.2851 0.1936 0.1754 0.1630 0.1829
## Balanced Accuracy 0.9989   0.9962   0.9941   0.9931   0.9971
```

The accuracy of this model is 99.4%, so the out-of-sample is only 0.6% and hence, this random forest model highly improves the before CART-tree model.

## Predictions

We are going to predict the classe-label for test data with the random forest model that we have previously trained.

The first thing is to drop all the variables that aren't used:

```
test_new <- test[,-index]
test_new <- test_new[,-c(1,2,3,4,5,6,7)]
```

And lastly, we make the prediction we have been asked for:

```
pred <- predict(model_rf, test_new, type = "class")
print(as.character(pred))

## [1] "B" "A" "B" "A" "A" "E" "D" "B" "A" "A" "B" "C" "B" "A" "E" "E" "A"
## [18] "B" "B" "B"
```